

TIDE: Indexing Time Intervals by Duration and Endpoint


Kai Wang, Moin Hussain Moti, Dimitris Papadias

Hong Kong University of Science and Technology


Aug 27, 2025

Interval Queries

- In temporal databases each record has a lifespan $[t_s, t_e)$, where $t_s < t_e$. A record is *alive* if t_e equals the current time; otherwise, it is *dead*.
- Stabbing query: returns intervals intersecting timestamp q .




- Range query: returns intervals intersecting range $[q_s, q_e]$.



Interval Tree [Edelsbrunner(1980)]


- Construct a BST (Binary Search Tree) by sorting endpoints of intervals.
- Pick the *median* endpoint as the root node recursively.
- An interval is stored at the *highest* node that intersects it.
 - E.g., interval $s_7 = [2, 22)$ is stored in n_{14} .



- Secondary structure for each node: two lists of intervals sorted on the two endpoints.


Segment Tree [De Berg(2000)]

- At the bottom level, each box is a conceptual leaf.
- Each interval is partitioned and stored in *multiple* nodes.
 - An interval is first stored at leaf nodes that intersect it.
 - Consecutive partitions are merged at the upper level recursively, if they cover their parent node.
 - E.g., $s_7 = [2, 22]$ is stored in n_4 , n_{12} , n_{18} and n_{23}^L .




Diagonal Corner Structure [Kanellakis et al.(1993)]

- Represent intervals as points in a 2D space.
- t_s is the horizontal axis, and t_e is the vertical axis.




stabbing query




range query

Start/End timestamp B-tree (SEB) [Song and Roussopoulos(2003)]

- A disk-based corner index
- Increasing Ending Time (IET) assumption:
intervals are inserted in increasing order of t_e .



Corner structure




index

- Two layers of *append-only* B+-trees:
 - Top tree: ordering intervals by t_s
 - Bottom tree: ordering intervals by t_e
- Append-only insertions: at the last data node of some bottom tree.

Queries in SEB

- SEB must search all bottom trees with $t_s \leq q_e$.
(Any interval that starts before q_e may die after q_s .)




A Unified Representation for Interval Indexes


- Interval indexes can be captured by some corner structure in a 2D space defined by t_s , t_e , or duration d .
- Benefit for regular queries
 - Identify nodes that must contain query results (i.e., all intervals can be directly reported).
- Benefit for *aggregate* queries
 - To compute the count of intervals intersecting a range, the number of intervals within each node inside the range can be aggregated directly, without visiting the node.

Interval Tree in the Corner Structure

- Each node is mapped to a rectangular area. E.g.,
 - n_{14} is mapped to the space of $t_s \leq 14 \leq t_e$.
 - n_6 is mapped to the space of $t_s \leq 6 \leq t_e < 14$.




the mapped structure



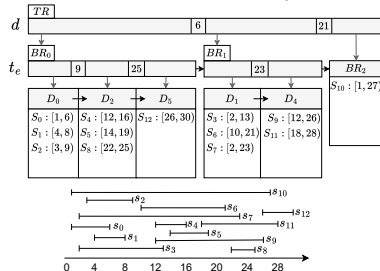
query with range $[9, 13]$

TIDE (time intervals by duration and endpoint)

- IET assumption: interval arrive in increasing order of t_e .
(i.e., they are inserted in the database when they die)



corner structure





index

- Two layers of *append-only* B+-trees:
 - Top tree: ordering intervals by duration d
 - Bottom tree: ordering intervals by t_e
- Append-only insertions: at the last data node of some bottom tree.



Insertions of TIDE

- Insert s_0, s_1, s_2 (assuming data node capacity 3)




Insertions of TIDE

- Insert s_3 : vertical split





Insertions of TIDE

- Insert s_4 : horizontal split





Insertions of TIDE

- Insert s_5, s_6, s_7, s_8 (according to the duration)




Datasets

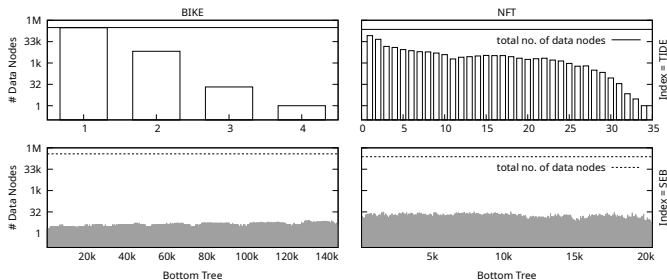
- **BIKE**: Start and end timestamps of 100M bicycle trips, during 2014-2020 in New York City.
- **NFT**: 28M intervals denoting the stable price period of non-fungible tokens transactions, during 2021-2023.




- I/O cost of sequential insertions (4 MiB buffer)



- Data nodes per bottom tree




- *Range queries* return qualifying records



- SEB searches numerous bottom trees with $t_s \leq q_e$.
 - For instance, a stabbing query in the middle of the data space is expected to visit at least half of the bottom trees.
- On BIKE, bottom trees with small t_s rarely contain results, although their mutable nodes intersect q_s .
- On NFT, even a stabbing query retrieves 1 million intervals, and this number remains almost the same for all ranges.

- *Count queries* return the number of qualifying records



- Compared to range queries, both TIDE and SEB incur less I/O. (aggregate directly results of full nodes covered by the query)
- TIDE has shorter nodes that may be covered by the query.
- SEB examines numerous irrelevant mutable nodes, which cannot be covered since they are open-ended.

Summary

- We propose a unified representation that
 - treats intervals as points in a 2D corner space
 - enables optimization opportunities for query processing
 - reveals strengths and weaknesses of interval indexes
- Following the IET assumption, TIDE outperforms SEB in
 - compactness and cache locality
(i.e., most non-full nodes are cached)
 - insertion efficiency (up to x100)
 - query efficiency (up to x7000)

References I



Mark De Berg. 2000.

Computational geometry: algorithms and applications.
Springer Science & Business Media.



H Edelsbrunner. 1980.

Dynamic rectangle intersection searching.
Technical Report (1980), 47.



Paris C. Kanellakis, Sridhar Ramaswamy, Darren E. Vengroff, and Jeffrey S. Vitter. 1993.

Indexing for data models with constraints and classes (extended abstract). In *Proceedings of the Twelfth ACM SIGACT-SIGMOD-SIGART Symposium on Principles of Database Systems* (Washington, D.C., USA) (PODS '93). Association for Computing Machinery, New York, NY, USA, 233–243.
[doi:10.1145/153850.153884](https://doi.org/10.1145/153850.153884)



Zhexuan Song and Nick Roussopoulos. 2003.

SEB-tree: An Approach to Index Continuously Moving Objects. In *Proceedings of the 4th International Conference on Mobile Data Management (MDM '03)*. Springer-Verlag, Berlin, Heidelberg, 340–344.